

Data Analytics Specialist Test

Cleaning, Validation, and Analysis of a Historical Aviation Crashes

Jorge A. Garcia

February 23, 2026

Outline

- 1 About me
- 2 General approach
- 3 Project structure
- 4 Data Cleaning
- 5 Data Validation
- 6 Analysis
- 7 Summary and Improvements

About me

- Quantitative modeler with 10+ years of experience:
 - Optimization (e.g., LP, ML)
 - Agent-based simulation
 - Physics-based modeling
 - Macroeconomic models
- Currently, developer of web-apps for science and decision-making (<https://modelbridgelabs.com/>).
- Goals:
 - Apply data analytics and modeling to inform decision-making and uncover actionable insights.
 - Learn and grow as a professional by taking on new challenges and responsibilities.

General approach

- ① Exploratory analysis
- ② Structure design
- ③ Data cleaning
- ④ Automated validation
- ⑤ Data profiling
- ⑥ Version control, CI, testing
- ⑦ Iterative refinement

```
. └── analytics
      └── analysis.py
          └── validation.py
    └── main.py
    └── rawdata
        └── plane_crashes_data.db
            └── plane_crashes_field_descriptions.csv
    └── README.md
    └── test.py
    └── utilspc
        └── cleanerclass.py
            └── cleanfun.py
```

Figure: Project structure.

Project structure

- `utilspc/`
 - `cleanfun.py` — Parsing functions for cleaning specific fields.
 - `cleanerclass.py` — Python classes for SQLite connection and cleaning.
- `analytics/`
 - `validation.py` — validation checks and profiling report generation.
 - `analysis.py` — analysis functions for generating reports and insights.
- `main.py` — main script to run the cleaning, validation, and analysis pipeline.
- `test.py` — test for type checking, cleaning functions, and validation checks.
- `output`
 - `metadata.csv` — column-level profile of the cleaned dataset.
 - `data_profile.txt` — full analysis report of the cleaned dataset.
 - `cleaned_plane_crashes.db` — cleaned SQLite database.
 - `validation_report.txt` — detailed validation report with check results and failure details.

Date, Time and Location

- **date**: Parse 'DD-Mon-YY' to 'YYYY-MM-DD', or None/Null on failure. Year values ≥ 26 are treated as 2000s, otherwise 1900s.
- **time**: Normalise to 'HH:MM' (24-hour). Strips prefix 'c' with optional space. Handles 4-digit integers (e.g. '1730' to '17:30'). Returns None for unparseable values.
- **location** Returns a tuple ('first', 'last') where 'first' is the location excluding the country (or None) and 'last' is the country name (or None). Creates field **country**.

Operator, AC type, aboard and fatalities

- **operator:** Strip whitespace, map ‘?’ to NULL and nullify values that are clearly not airline operators:
 - Pure serial / registration codes, e.g. “46826/109”.
 - Aircraft manufacturer + model designator entries that were incorrectly placed in the operator column, e.g. ‘Boeing KC-135E’, ‘Lockheed AC-130H Hercules’.
- **ac_type:**
 - Strip whitespace and map ‘?’ to NULL.
 - Remove vehicle categories, e.g. ‘(flying boat)’, ‘(airship)’, ‘(amphibian)’.
 - Remove extra spaces.
 - Strip any residual leading ‘/’ trailing whitespace.
 - Return None if empty
- **aboard, fatalities:** Returns (total, passengers, crew) as int or None.

Data Validation

Automated checks run after every cleaning pass:

Structure

- Schema: all 18 expected columns with correct declared types
- Python-level type consistency per column

Format

- Dates match YYYY-MM-DD; range 1908–2018
- Times match HH:MM (24-hour)
- All numeric columns ≥ 0

Cross-column consistency

- $\text{aboard_pax} + \text{aboard_crew} = \text{aboard_total}$
- $\text{fat_pax} + \text{fat_crew} = \text{fatalities_aboard}$
- $\text{fatalities_aboard} \leq \text{aboard_total}$
- $\text{fatalities_total} = \text{fatalities_aboard} + \text{ground}$

Duplicates

- Rows sharing (date, operator, route) flagged as potential duplicates

Results reported as **PASS / WARN / FAIL** per check.

Analysis

Five sections produced in the profiling report:

- ① **Data Profiling** — NULL rate, unique value count, and type per column.
- ② **Descriptive Statistics** — Aggregate fatality rate (fatalities / aboard), survival rate, ground casualties, and crew vs. passenger fatality split.
- ③ **Trend Analysis** — Crashes and fatalities per decade and per year (ASCII bar chart); top 15 operators by crash count; most dangerous aircraft types by fatality rate (min. 10 incidents).
- ④ **Geographic Analysis** — Top 20 countries/regions and specific crash sites by incident count.
- ⑤ **Data Quality** — Mismatched totals, rows where fatalities exceed aboard count, duplicate (date, operator, route) groups, and registration reuse across aircraft types.

Summary and Future Work

My approach:

- Develop a modular and reusable structure.
- Implement programming best practices: version control, testing, CI/CD, and documentation.
- Explore ways to integrate data products with decision-making processes and end-user applications.

Improvements:

- Enhance location parsing to better handle edge cases and extract more granular geographic information.
- Implement additional validation checks for outliers and temporal inconsistencies (e.g., crashes dated before the Wright brothers' first flight).
- Standardize operator names, AC types, and geographic scales.
- Complement data with external reference datasets (e.g., GIS, weather, economic indicators).