

Data Analytics Specialist Test

Cleaning, Validation, and Analysis of a Historical Aviation Crashes

Jorge A. Garcia

February 23, 2026

Outline

- 1 About me
- 2 General approach
- 3 Project structure
- 4 Data Cleaning
- 5 Data Validation
- 6 Metadata
- 7 Analysis
- 8 User Interface
- 9 Summary
- 10 Disclosure

About me

- Quantitative modeler with 10+ years of experience:
 - Optimization (e.g., LP, ML)
 - Agent-based simulation
 - Physics-based modeling
 - Macroeconomic models
- Currently, developer of web-apps for science and decision-making (<https://modelbridgelabs.com/>).
- Goals:
 - Apply data analytics and modeling to inform decision-making and uncover actionable insights.
 - Learn and grow as a professional by taking on new challenges and responsibilities.

General approach

- ① Exploratory field analysis
- ② Structure design
- ③ Data cleaning
- ④ Automated validation
- ⑤ Data profiling
- ⑥ Version control, CI, testing
(github.com/jorge-antares/das_test)
- ⑦ Iterative refinement

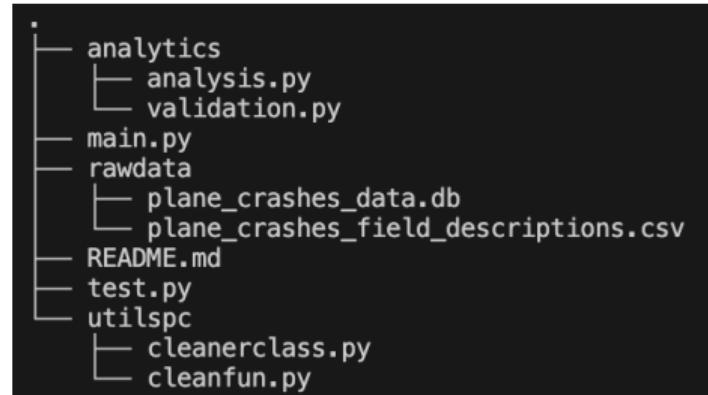


Figure: Project structure.

Project structure

- `utilspc/`
 - `cleanfun.py` — Parsing functions for cleaning specific fields.
 - `cleanerclass.py` — Python classes for SQLite connection and cleaning.
- `analytics/`
 - `validation.py` — validation checks and profiling report generation.
 - `analysis.py` — analysis functions for generating reports and insights.
- `main.py` — main script to run the cleaning, validation, and analysis pipeline.
- `test.py` — simple validity test (CI - GitHub Actions).
- `output`
 - `metadata.csv` — column-level profile of the cleaned dataset.
 - `data_profile.txt` — full analysis report of the cleaned dataset.
 - `cleaned_plane_crashes.db` — cleaned SQLite database.
 - `validation_report.txt` — detailed validation report with check results and failure details.

Date, Time and Location

- **date**: Ambiguity in YY-MM-DD or DD-MM-YY formats is resolved by assuming the former. Parse 'DD-Mon-YY' to 'YYYY-MM-DD', or None/Null on failure. UTF?
- **time**: Normalise to 'HH:MM' (24-hour). Strips prefix 'c' with optional space. Handles 4-digit integers (e.g. '1730' to '17:30'). Returns None for unparseable values.
- **location** Returns a tuple ('first', 'last') where 'first' is the location excluding the country (or None) and 'last' is the country name (or None). Creates field **country**.

Operator, AC type, aboard and fatalities

- **operator:** Strip whitespace, map ‘?’ to NULL and nullify values that are clearly not airline operators:
 - Pure serial / registration codes, e.g. “46826/109”.
 - Aircraft manufacturer + model designator entries that were incorrectly placed in the operator column, e.g. ‘Boeing KC-135E’, ‘Lockheed AC-130H Hercules’.
- **ac_type:**
 - Strip whitespace and map ‘?’ to NULL.
 - Remove vehicle categories, e.g. ‘(flying boat)’, ‘(airship)’, ‘(amphibian)’.
 - Remove extra spaces.
 - Strip any residual leading ‘/’ trailing whitespace.
 - Return None if empty
- **aboard, fatalities:** Returns (total, passengers, crew) as int or None.

Data Validation

Automated checks run after every cleaning pass:

Structure

- Schema: all 18 expected columns with correct declared types
- Python-level type consistency per column

Format

- Dates match YYYY-MM-DD; range 1908–2025
- Times match HH:MM (24-hour)
- All numeric columns ≥ 0

Cross-column consistency

- $\text{aboard_passengers} + \text{aboard_crew} = \text{aboard_total}$
- $\text{fatalities_passengers} + \text{fatalities_crew} = \text{fatalities_aboard}$
- $\text{fatalities_aboard} \leq \text{aboard_total}$
- $\text{fatalities_total} = \text{fatalities_aboard} + \text{ground}$

Duplicates

- Rows sharing (date, operator, route) flagged as potential duplicates

Results reported as **PASS / WARN / FAIL** per check.

Dataset Schema

metadata.csv (5,783 rows)

Field	Type	NA	Unique	Description
date	TEXT	0	5212	Date of the crash (YYYY-MM-DD)
time	TEXT	2183	1051	Time of the crash (HH:MM, 24-hour)
location	TEXT	230	4372	Location of the crash
country	TEXT	104	429	Country of crash (derived from location)
operator	TEXT	29	2812	Airline or operator
flight_no	TEXT	4508	868	Flight number assigned by the operator
route	TEXT	1689	3839	Route flown prior to the accident
ac_type	TEXT	24	2666	Aircraft type
registration	TEXT	357	5385	ICAO registration of the aircraft
cn_ln	TEXT	1209	4063	Construction/serial number
aboard_total	INTEGER	40	242	Total number of people aboard
aboard_passengers	INTEGER	543	231	Number of passengers aboard
aboard_crew	INTEGER	539	31	Number of crew aboard
fatalities_aboard	INTEGER	11	200	Total number of fatalities aboard
fatalities_passengers	INTEGER	558	190	Number of passenger fatalities
fatalities_crew	INTEGER	556	28	Number of crew fatalities
ground	INTEGER	52	53	Ground fatalities
fatalities_total	INTEGER	0	204	Total number of fatalities
summary	TEXT	231	5366	Brief description of the accident

Analysis

Five sections produced in the profiling report:

- ① **Data Profiling** — NULL rate, unique value count, and type per column.
- ② **Descriptive Statistics** — Aggregate fatality rate (fatalities / aboard), survival rate, ground casualties, and crew vs. passenger fatality split.
- ③ **Trend Analysis** — Crashes and fatalities per decade and per year (ASCII bar chart); top 15 operators by crash count; most dangerous aircraft types by fatality rate (min. 10 incidents).
- ④ **Geographic Analysis** — Top 20 countries/regions and specific crash sites by incident count.
- ⑤ **Data Quality** — Mismatched totals, rows where fatalities exceed aboard count, duplicate (date, operator, route) groups, and registration reuse across aircraft types.

User Interface




Figure: Web app interface at
<https://40.233.98.115/>.

Frameworks used:

- Flask for backend API.
- Jinja2 for HTML templating.
- Bootstrap for responsive design.
- Docker for containerisation and deployment.
- Nginx for reverse proxy and static file serving.

Hosting via a VM instance (Ubuntu server 24.04.4 LTS) on Oracle Cloud.

Summary

My approach:

- Develop a modular and reusable structure with minimal external dependencies.
- Implement programming best practices: version control, testing, CI/CD, and documentation.
- Standardise fields to keep consistency and enable analysis.
- Cross-validate data and create metrics (e.g., fatality rates).
- Explore ways to integrate data products with decision-making processes and end-user applications.

Disclosure of Tools

Claude Sonnet 4.6 was used in the following tasks:

- Parse fields to look for edge cases.
- Create documentation based on project structure and code comments.
- Create HTML for user interface.
- Code completion and debugging suggestions (VS Code).

All LLM outputs were reviewed and edited.